

Predicting FEMA Hazard Mitigation Costs and State Trends

Anjana Dhakal

Northwest Missouri State University, Maryville MO 64468, USA
SS574360@nwmissouri.edu and dhakaal.anjanaa@gmail.com

Abstract. This capstone project analyzes the FEMA Hazard Mitigation Assistance (HMA) Mitigated Properties dataset (v4) to uncover patterns in disaster mitigation spending and predict future project costs. Focusing on U.S. states post-2000, the study employs exploratory data analysis (EDA) to reveal skewed cost distributions, reactive funding trends, and geographic hotspots. Supervised regression models-Linear Regression (LR) and Random Forest Regressor (RFR)-predict log-transformed costs, with RFR selected for superior performance (test $R^2 = 0.5847$ on log scale). Projections for FY2030 indicate a 17.9% increase in national spending (\$1.63B to \$1.92B), but sharp declines in historical hotspots, signaling the risks of risk redistribution. Insights emphasize proactive, equitable funding to enhance resilience amid rising disasters.

Keywords: FEMA· Hazard Mitigation· Project Cost Analysis· Exploratory Data Analysis (EDA)· Linear Regression Forecasting· Random Forest Forecasting

1 Introduction:

1.1 Background

Disaster mitigation is an essential element of national resilience planning, particularly as the frequency and economic impact of natural disasters continue to increase in the United States (Borden et al., 2007). The Federal Emergency Management Agency (FEMA) plays a key role in managing hazard mitigation programs aimed at reducing future disaster losses. Through its Hazard Mitigation Assistance (HMA) programs, FEMA funds projects that improve community infrastructure, prevent repeated damage, and lower long-term recovery costs through measures such as property acquisition, flood control, and infrastructure hardening (FEMA, 2025; Dixon & Johnson, 2019).

Studies have shown that hazard mitigation efforts significantly reduce disaster related losses. For example, (Rose et al., 2007) demonstrated that every dollar invested in FEMA’s mitigation grants yields multiple dollars in avoided future disaster costs. Their benefit-cost analysis found that FEMA-funded projects are cost-effective and socially beneficial regardless of the type of danger.

1.2 Significance

The project analyzes how FEMA Hazard Mitigation Assistance (HMA) project costs vary across US states post 2000, hazard types, and property characteristics. Identifying hotspots and predictors enables more targeted allocation of limited federal funds, enhancing preparedness in vulnerable regions. This problem is also technically valuable

as there are not many studies that combine exploratory analysis and predictive modeling with visualization.

1.3 Goals of this Project

The primary objectives are:

- Conduct an EDA on post-2000 HMA data, identifying cost trends by fiscal year, state, and project characteristics.
- Build and evaluate regression models (Linear Regression and Random Forest) to predict project-level mitigation costs.
- Visualize patterns and recommend strategies for equitable funding.

2 Data Description:

2.1 Data Source, Format, and Size

The FEMA Hazard Mitigation Assistance (HMA) dataset was obtained from the FEMA Open Data Portal. It is provided in CSV format, approximately 15 MB in size, containing 97,538 records and 19 columns in its original form. For this study, the dataset was filtered to include only U.S. states and projects completed after fiscal year 2000, resulting in 84,026 records across 11 key variables.

Data preparation steps standardized categorical variables, handled missing fields, and derived new attributes such as *totalMitigationCost*. These transformations ensured a consistent, analysis-ready dataset suitable for exploratory and predictive modeling.

2.2 Data Extraction

The FEMA Hazard Mitigation Assistance (HMA) dataset (v4) was obtained from the FEMA Open Data Portal. The original data source is available at [\(FEMA HMA Mitigated Properties Dataset\)](#).

Python tools such as Pandas and NumPy were used for efficient data manipulation, while Matplotlib and Seaborn supported visualization. All work was performed in VS Code using Pandas, ensuring a reproducible and transparent workflow from extraction through modeling.

2.3 Data Dictionary

The analysis leverages a subset of fields for focused exploratory data analysis (EDA) and modeling. Table 1 summarizes the selected attributes used in this project.

2.4 Data Quality and Cleaning

The dataset is fairly structured but contains known issues like missing values in optional fields, duplicate records, and invalid state entries. For this reason, an additional data cleaning process was carried out, which is explained in detail in Section 3.2 Data Cleaning.

3 Methodology

The methodology describes the step-by-step approach followed in this project, covering all stages from data collection to model evaluation.

Table 1. Selected fields from the FEMA HMA Mitigated Properties dataset. The table lists the attribute name, description, data type, and an example value for each field.

Field Name	Description	Data Type	Example Value
state	U.S. state where the property is located	String	Texas
county	County name where the property is located	String	Harris
damageCategory	Severity of damage (e.g., Substantial, Minor)	String	Substantial
structureType	Type of property structure (e.g., Single Family)	String	Single Family
typeOfResidency	Residency type (e.g., Owner-Occupied)	String	Owner-Occupied
foundationType	Property foundation type (e.g., Slab-On-Grade)	String	Slab-On-Grade
actualAmountPaid	Amount paid to property owner (USD; excludes administrative costs)	Float	150000.00
numberOfProperties	Number of properties mitigated in the project	Integer	5
propertyAction	Mitigation action type (e.g., Elevation, Acquisition)	String	Elevation
propertyPartOfProject	Indicates if property is part of a larger project	String	Yes
programFy	Fiscal year of the mitigation program	Integer	2024
date_approved_proxy	Approval date proxy derived from fiscal year	Datetime	2024-07-01
totalMitigationCost	Total mitigation cost (actualAmountPaid × numberOfProperties)	Float	750000.00



Fig. 1. Project implementation workflow,(A. Dhakal, 2025)

3.1 Data Collection

The dataset was obtained as a CSV file from the official FEMA Open Data website ([FEMA HMA Mitigated Properties Dataset](#)).

3.2 Data Cleaning

The data cleaning process transformed the raw FEMA HMA dataset (97,515 rows $\times$ 19 columns) into an analysis-ready format, addressing missing values, inconsistencies, duplicates, and irrelevancies. No joins were needed (single-source CSV), but derived features aggregated costs for modeling. Screenshots of pre/post summaries are in Appendix A.

The data cleaning steps included:

1. Early filter to US states and post-2000.
2. Keep relevant columns (including model predictors).
3. Handle missing values early (for categoricals before standardization).
4. Type conversions and standardize categoricals (with str conversion for safety).
5. Derive date proxy and clip negatives.
6. Remove duplicates and derive total cost.
7. Validation and save cleaned data in csv format.

Missing values:

Missing data were handled through targeted imputation and selective deletion. For categorical fields such as `damageCategory` (61% nulls), an `Unknown` label preserved 84,026 records. Numerical fields were treated as follows: `numberOfProperties` was imputed with 1, reflecting mostly single-property projects, and rows missing `actualAmountPaid` (48% nulls) were dropped, retaining 47,199 valid cases. These steps align with FEMA data standards and were validated through distribution checks to prevent bias in EDA and modeling.

Final Dataset

The final dataset is shown below:

Table 2. Summary of Data Cleaning Stages

Stage	Rows	Columns	Key Action
Raw	97,515	19	Initial load
Filtered	84,026	19	US states / post-2000 subset
Cleaned	41,308	13	Imputed, derived, and deduplicated

Attribute Definitions and Transformations:

Subset of Table 1, standardized (title-case; mappings applied, e.g., “N/A” → “Unknown”). New fields derived during cleaning:

- `date_approved_proxy`: *datetime* derived from fiscal year.
- `totalMitigationCost`: *float*, calculated as `actualAmountPaid × numberOfProperties`.

4 Exploratory Data Analysis

The EDA process summarized numeric distributions using histograms, boxplots, and a correlation heatmap, revealing strong right-skewness in *actualAmountPaid*, which was addressed with a `log1p` transformation. Temporal patterns were examined through fiscal-year aggregations, while mitigation strategies were explored using value counts, pie charts, and bar charts. Geographic concentrations were analyzed through state-level summaries and horizontal bar plots. All visualizations were produced with Seaborn and Matplotlib and saved as PNG files. Overall, the EDA highlighted three core insights: concentrated financial impacts, reactive temporal spending patterns, and FEMA’s broader mitigation strategy.

Log Transformation Code

```
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# Apply log1p to handle skewness and zeros
df['log_actualAmountPaid'] = np.log1p(df['actualAmountPaid'])

# Normalized view
sns.histplot(df['log_actualAmountPaid'], kde=True)
plt.title('Log-Transformed Cost Distribution')
```

This transformation reveals underlying patterns hidden by outliers, enabling more reliable visualizations and modeling.

5 Modeling and Analysis

The analysis used supervised regression to predict *actualAmountPaid* (log-transformed) based on four categorical features (*state*, *structureType*, *typeOfResidency*, *foundationType*) and one numeric feature (*numberOfProperties*). A `ColumnTransformer` with

OneHotEncoder handled preprocessing, and the data was split 80/20 for training and testing. Two models were evaluated: a baseline Linear Regression and a Random Forest Regressor designed to capture non-linear interactions in the one-hot encoded feature space. Performance was assessed using R^2 on the log scale and MAE/RMSE on the original scale. The model with the higher log-scale R^2 was selected for state-level projections to estimate FY2030 outcomes, supported by bar and line visualizations.

Code Example: Model Training and Evaluation

Random Forest Model Training

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score

rf_model = RandomForestRegressor(
    n_estimators=100,
    max_depth=15,
    random_state=42
)

rf_model.fit(X_train, y_train) # y_train = log_costs
y_pred_log = rf_model.predict(X_test)

r2_log = r2_score(y_test, y_pred_log)
# High R2 indicates strong fit
print(f"Random Forest R2 (Log): {r2_log:.4f}")
```

Random Forest's ensemble approach captures non-linear cost drivers, outperforming linear baselines for accurate projections.

6 Results

All results presented here are based on the Python notebooks for this project, with all code available at: [[GitHub Notebook Link](#)].

Disasters are increasing, and FEMA's HMA program provides protection, but spending is uneven-focused on high-risk areas, primarily homes, and surging after major events. The analysis below examines how costs, strategies, timing, and geography shape national resilience.

6.1 Uncovering Extreme Cost Concentration and Model Preparation

The distribution of money paid by FEMA reveals the impacts of major disasters and their associated costs driving the costs higher by those disasters. The average project cost (\$272, 325) is four times the 75th percentile (\$66, 046) (Table 3). This extreme concentration validates the focus of the project: a small fraction of high-cost projects drives the entire mitigation budget. Although majority of the projects cost 0 to 100 million dollars, there are a few which are significantly expensive with the cost almost a billion dollars (Fig.2). Even the cheaper projects are not cheap as the graph in fig.2

shows the values in hundred millions (10^8). Most of the projects cost one hundred to two hundred million dollars.

Table 3. Initial descriptive statistics for the pre-modeling dataset. Shows count, mean, and quartiles for key variables including Actual Amount Paid, Number of Properties, and Total Mitigation Cost.

Variable	Count	Mean	25%	75%
Actual Amount Paid (\$)	41,308	272,325	3,972	66,046
Number of Properties	41,308	1.58	1	1
Total Mitigation Cost (\$)	41,308	299,917	4,703	69,500




Fig. 2. Distribution of Actual Amount Paid, Number of Properties, and Total Mitigation Cost. Histograms and boxplots show strong right-skewness across all three variables, with most projects concentrated at lower cost ranges and a long tail extending into multi-million-dollar amounts. The presence of substantial outliers further supports applying a target variable transformation. (A. Dhakal, 2025)

Correlation Analysis:

A correlation heatmap was generated to inspect the relationships between the numeric features before modeling. Fig 3. Correlation Heatmap of Numeric Variables suggests that the actual amount paid is strongly correlated to totalmitigation cost but it is not strongly correlated to number of properties. It could be because of the fact that there are fewer numbers of properties that cost a lot and majority of properties do not cost high amounts in different states.




Fig. 3. Correlation heatmap of numeric variables, showing strong positive correlation between Actual Amount Paid and Total Mitigation Cost and low correlation with Number of Properties. (A. Dhakal, 2025)

Target Transformation:

The target variable exhibited extreme right-skewness, a normalization step was required for robust regression modeling. A $\log(x + 1)$ transformation was applied to stabilize the variance and reduce the influence of outliers.

Code Example (Transformation):

Log Transformation

```
import numpy as np
# Applying log(x + 1) transformation
df['log_cost'] = np.log1p(df['actualAmountPaid'])
```

The transformation successfully converted the highly skewed target into an approximately normal distribution, enabling more stable and reliable model training.




Fig. 4. Target variable distribution before and after $\log(x + 1)$ transformation. Log transformation reduces skewness and stabilizes variance for regression modeling.(A. Dhakal, 2025)

6.2 Mitigation Strategy: Dual Approaches to Risk

Frequency analysis was used to determine FEMA’s strategic priorities by propertyAction and structureType. Most of the mitigation funding is concentrated on certain types of risk-reduction strategies. Wind Retrofit and Acquisition/Demolition account for more than 70% of all spending. Most of the mitigation budgets are spent in supporting Single Family structures.



Fig. 5. Distribution of mitigation actions by propertyAction and structureType. Wind retrofits and acquisition/demolition dominate, with single-family homes receiving most funding.(A. Dhakal, 2025)

6.3 Funding Trends: Peaks After Disaster

Federal funding by FEMA spikes in response to major disaster. From Fig.6, the fiscal years 2008, 2011, and 2017 show a high increase in both total dollars paid and properties mitigated. These years correspond to major disasters in recent U.S. history, including post-Hurricane Katrina recovery and Midwest floods (2008), the Joplin tornado and Mississippi River floods (2011), and Hurricanes Harvey, Irma, and Maria (2017).

Code Example:

Yearly Aggregation

```
yearly_trends = df.groupby('programFy').agg({
    'numberOfProperties': 'sum',
    'actualAmountPaid': 'sum'
}).reset_index()
```




Fig. 6. Yearly funding trends showing reactive spikes after major disasters. Peaks follow post-disaster years, indicating FEMA's mitigation is largely event-driven. (A. Dhakal, 2025)

6.4 Geographic Concentration: High-Cost Hotspots

All of the states receive one form or another mitigation dollars based on the disasters they encounter. Among them Missouri (\$2.6B), Washington (\$1.7B), and Florida (\$1.1B) account for 45 percent of all dollars spent (Fig. 7). This result suggests targeted investment based on the repeated disasters they encounter. These investments also raise questions about balancing urgent need in some particular areas in response to a disaster versus developing capacity for the entire nation.

These hotspots reflect regional hazards:

- Missouri's large-scale flood buyouts
- Washington's seismic retrofits
- Florida's hurricane-hardening efforts




Fig. 7. Top 10 states by total mitigation spending. Missouri, Washington, and Florida receive the largest share, reflecting regional hazard exposure.(A. Dhakal, 2025)

6.5 Modeling and Analysis

6.5.1 Predictive Model Setup and Comparison

To address the wide range in project costs, the target variable was transformed using $\log(\text{actualAmountPaid} + 1)$. This transformed outcome was modeled using key features including state, structureType, foundationType, numberOfProperties, and programFy. Categorical variables were one-hot encoded, and numerical variables were used as-is. Two predictive models-Linear Regression and Random Forest Regressor-were trained and compared using the R^2 score on the log-transformed test data. An 80/20 train-test split (random state = 42) ensured consistent evaluation. Random Forest achieved higher R^2 values for both training and test sets (Table 4). R^2 values near zero on the original scale do not indicate poor performance. They simply reflect the extreme right-skew of the cost data. The log transformation is necessary to stabilize variance, reduce outlier influence, and provide a more meaningful basis for evaluation.

Table 4. Model performance comparison for LR and RFR. Metrics include R^2 , MAE, and RMSE on training and test datasets (log-transformed and original).

Dataset	Model	R^2 (Log)	R^2 (Original)	MAE (\$)	RMSE (\$)
Training	Linear Regression	0.3955	-0.0006	265,200	9,092,721
Training	Random Forest Regressor	0.6570	0.0019	253,640	9,081,705
Test	Linear Regression	0.3857	-0.0003	222,836	9,635,366
Test	Random Forest Regressor	0.5847	-0.0003	213,165	9,634,923

Table 5. Five-fold cross-validation summary for LR and RFR. Shows mean and standard deviation of R^2 on log-transformed data and test R^2 .

Model	Mean CV R^2 (Log)	Std CV R^2 (Log)	Test R^2 (Log)
Linear Regression	0.3858	0.0112	0.3857
Random Forest Regressor	0.5032	0.0105	0.5847

6.5.2 Factors Impacting the Mitigation Costs

Program Fiscal Year was the strongest predictor (0.13), indicating mitigation costs trend upward over time. The State of Oklahoma followed (0.12), likely reflecting tornado-related disasters. Foundation types and the number of properties contributed less.

6.5.3 Top 5 State Projection (RFR Model)

The Random Forest Regressor predicted mitigation costs for the full dataset and for a projected future scenario where *programFy* was increased by five years (simulating FY2030). Predictions were inverse-transformed to USD and aggregated by state. National costs rose 17.9% (\$1.63B $\rightarrow$ \$1.92B). However, linear trends for the top 5 states showed sharp declines (Table 6, Fig. 8), possibly reflecting lower recent spending compared with major historical events.

Table 6. Top 5 states by historical and projected FY2030 mitigation spending (RFR). Shows declines in historically high-spending states, indicating saturation effects.

State	Current Cost (\$M)	Projected FY2030 Cost (\$M)
Missouri	2,605	240
Washington	1,703	19
Florida	1,082	251
Alabama	1,025	107
Pennsylvania	1,006	62




Fig. 8. Current versus projected FY2030 mitigation spending for the top 5 states. Declines in historically high-spending states suggest saturation effects and potential funding shifts.(A. Dhakal, 2025)

7 Conclusions and Discussion

This analysis of more than 41,000 FEMA Hazard Mitigation Assistance (HMA) projects from 2000-2025 shows that almost 1.63 billion USD is spent in responding to hazards and on mitigation costs in the U.S. A majority of such projects are small but there are a few huge projects which skew the results. Majority of such budgets seem to be spent on a few categories like wind retrofits and property acquisitions/demolitions. These two categories account 70% of all activities (Fig. 5). One highlights of thee result is a focus on single-family homes. This suggests that the projects help working family to be more resilient and recover from the disasters.

Funding trends highlight the reactive nature of the FEMA's assistance. The large spending correspond to the major disaster periods like 2008, 2011, and 2017. These periods relate to post-disaster periods following Hurricanes Katrina, Sandy, and Harvey/Irma (Fig. 6). The geographic pattern adds another imbalance: Missouri (\$2.6B), Washington (\$1.7B), and Florida (\$1.1B) account for 45% of total spending (Fig. 7), leaving many regions comparatively underserved. These disparities raise equity concerns as disaster risks grow nationwide.

The Random Forest Regressor (RFR) provides further insight by capturing non-linear relationships that linear regression misses (test $R^2 = 0.5847$ vs. 0.3857 ; Table 5). The model identifies fiscal year (0.1285) and state-specific factors (e.g., Oklahoma at 0.1222) as the strongest cost drivers, while project count plays a smaller role (0.0613). Cross-validation supports model stability (mean $R^2 = 0.5032$; Table 6). Forecasts suggest a 17.9% increase in national spending by FY2030 (from \$1.63B to \$1.92B), paired with sharp slowdowns in historically high-investment states: projected declines of 91% in Missouri and 99% in Washington (Table 7; Fig. 8). This pattern reflects saturation in high-spend areas and steady national growth, emphasizing FEMA's challenge of balancing episodic surges with sustained, equitable investment.

In summary, HMA effectively addresses immediate crises but falls short on proactive prevention, leaving communities vulnerable. These findings call for a shift toward

data-driven, equitable strategies that turn episodic relief into lasting resilience (Rose et al., 2007).

7.1 Limitations

Here are some of the limitations of this study:

1. It includes data reported, maintained, and made publicly available by FEMA only.
2. The FEMA dataset is heavily weighted toward residential projects (over 70%), which may underrepresent public or commercial infrastructure needs.
3. Data cleaning removed 57% of original records (from 97,515 to 41,308), largely due to extensive missing fields (e.g., 59,807 null values in `damageCategory`), potentially excluding lower-severity projects.
4. The RFR model also assumes temporal stability, which may not hold true as the disasters are rarely linear in occurrence.

7.2 Recommendations

1. **Shift to proactive funding:** Increase HMA support during non-disaster years to help low-funding states and reduce risks before crises occur.
2. **Diversify project focus:** Continue supporting residential work but expand investment in public and commercial infrastructure, which remains underfunded.
3. **Risk-based allocation:** Apply Random Forest insights-especially the strong influence of state and foundation type to design risk-indexed grants that account for future exposure and project complexity rather than relying solely on past losses.

8 References

1. Borden, K. A., Schmidtlein, M. C., Emrich, C. T., Piegorsch, W. W., & Cutter, S. L. (2007). *Vulnerability of U.S. Cities to Environmental Hazards*. *Journal of Homeland Security and Emergency Management*, 4(2).
2. Dixon, M. J., & Johnson, K. A. (2019). *Hazard Mitigation Policy and Practice in the United States*. In *The Oxford Research Encyclopedia of Natural Hazard Science*. Oxford University Press.
3. FEMA. (2025). *Hazard Mitigation Assistance Mitigated Properties - v4 (Open-FEMA Dataset)*.
4. Rose, A., et al. (2007). *Benefit-Cost Analysis of FEMA Hazard Mitigation Grants*. *Natural Hazards Review*, 8(4), 97–111.

9 Appendix

Supplementary Materials

The analysis, cleaned datasets, and code are available at the following locations:

1. Overleaf link :[Overleaf Project](#)
2. GitHub Repo link: [GitHub Repo](#)

Appendix A: Pre- and Post-Cleaning Dataset Summaries

This appendix provides textual representations of key pre- and post-cleaning summaries from the Jupyter Notebook outputs. These include dataset info, null value counts, and data types, generated via Pandas `.info()`, `.isnull().sum()`, and `.dtypes`.

A.1 Raw Dataset Summary (Pre-Cleaning)

Dataset Shape: 97,515 rows $\times$ 19 columns

Data Overview: The dataset contains 97,515 entries, indexed from 0 to 97,514.

Column Information and Data Types:

#	Column	Non-Null Count	Dtype
0	projectIdentifier	97515 (non-null)	obj
1	programArea	97515 (non-null)	obj
2	programFy	97515 (non-null)	int
3	disasterNumber	85184 (non-null)	flt
4	propertyPartOfProject	93361 (non-null)	flt
5	propertyAction	90837 (non-null)	obj
6	structureType	90427 (non-null)	obj
7	typeOfResidency	61502 (non-null)	obj
8	foundationType	87029 (non-null)	obj
9	county	94453 (non-null)	obj
10	city	97511 (non-null)	obj
11	state	97477 (non-null)	obj
12	stateNumberCode	97476 (non-null)	flt
13	region	97477 (non-null)	flt
14	zip	97505 (non-null)	obj
15	damageCategory	37708 (non-null)	obj
16	actualAmountPaid	50352 (non-null)	flt
17	numberOfProperties	97515 (non-null)	int
18	id	97515 (non-null)	obj

Null Value Counts (Sorted Descending):

Column	Null Count
damageCategory	59,807
actualAmountPaid	47,163
typeOfResidency	36,013
disasterNumber	12,331
foundationType	10,486
structureType	7,088
propertyAction	6,678
propertyPartOfProject	4,154
county	3,062
stateNumberCode	39
region	38
state	38
zip	10
city	4
projectIdentifier	0
programFy	0
programArea	0
numberOfProperties	0

```
id                                0
```

A.2 Cleaned Dataset Summary (Post-Cleaning)

Final Dataset Shape: 41,308 rows $\times$ 13 columns

Column Information and Data Types:

Column	Dtype
state	obj
county	obj
damageCategory	obj
structureType	obj
typeOfResidency	obj
foundationType	obj
actualAmountPaid	flt
numberOfProperties	int
propertyAction	obj
propertyPartOfProject	obj
programFy	int
date_approved_proxy	dt
totalMitigationCost	flt

Null Value Counts (All Zero):

Column	Null Count
state	0
county	0
damageCategory	0
structureType	0
typeOfResidency	0
foundationType	0
actualAmountPaid	0
numberOfProperties	0
propertyAction	0
propertyPartOfProject	0
programFy	0
date_approved_proxy	0
totalMitigationCost	0